

# Auditable Flourishing: A Non-Compensatory Method for Assuring Frameworks That Claim to Support Flourishing

**Methods article submission draft: Version 6.2**

James McGaughran

ORCID: 0009-0005-3324-7290

British University Vietnam

## Abstract

Decision frameworks increasingly claim to improve flourishing, wellbeing, sustainability, justice, safety, or aligned governance. Those claims often combine normative commitments, empirical models, institutional authority, measurement systems, and implementation rules. Existing evaluation and assurance traditions provide valuable parts of the needed method, but a recurring gap remains: before comparing frameworks or using their outputs, institutions need a disciplined way to determine whether the frameworks themselves are inspectable, rights-protecting, auditable, measurable, corrigible, and capable of refusal. Auditable Flourishing (AF) is a proposed non-compensatory assurance protocol for that purpose. It evaluates candidate decision architectures and their evidentiary claims; it does not directly measure or certify flourishing, establish legal or moral authority, or authorize deployment.

The method uses two stages. Stage A applies six non-compensatory eligibility criteria: C0 Inspectability, C1 Rights, C2 Auditability, C3 Measurement integrity, C4 Corrigibility, and C5 Refusal and execution boundary. A material defect cannot be averaged away by strength elsewhere. Stage A produces bounded, scope- and version-specific outcomes, including scoped admissibility, scoped control-dependent admissibility, patchable non-admissibility, unpatchable non-admissibility, indeterminacy, and protocol refusal. Stage B compares only Stage-A-admissible candidates under a common scope, evidence cut-off, frozen scenario set, and provisional twelve-dimension ordinal profile. Ordinal categories are not summed or converted into a pseudo-metric score. The protocol computes a partial-order relation and replays every evidence-supported category perturbation. Invariant one-direction dominance supports a dominance finding; invariant crossing supports incomparability; relation changes, ties, or unresolved but potentially repairable conditions produce non-decisiveness; invalid prerequisites produce refusal to compare.

Version 6.2 closes independent-audit findings and consolidates the operational safeguards required for independent testing: immutable pre-discussion rating deposits, randomized criterion order, global-impression probes, scenario-set governance, reviewer qualification, host-interest and AI-assistance disclosure, conjunctive burden and access feasibility, an authority-specific verdict vector, an anti-liability-shield rule, and a feasible exploratory Phase 0 study. The framework is internally specification-complete and Phase 0 design-ready, but inter-reviewer reliability, construct validity, architecture fairness, decision value, affected-party legitimacy, safe live administration, and comparative superiority remain open empirical questions. AF is therefore presented as a falsifiable research and assurance method, not a completed certification standard.

**Keywords:** assurance; framework evaluation; non-compensation; rights floors; auditability; ordinal profiles; partial orders; corrigibility; refusal; flourishing

## 1. The assurance problem

Public and private institutions increasingly rely on recurring decision frameworks rather than isolated judgments. A framework may be a legal or policy procedure, a multi-criteria method, a public-sector algorithm, a community deliberative process, a safety case, an AI governance architecture, or a hybrid human-machine system. Once adopted, it determines what evidence is visible, which people have standing, which harms count, how uncertainty is represented, what can be traded off, and who may authorize action. Its architecture therefore shapes the outcome before any particular option is scored.

Evaluation usually begins too late. Analysts may compare outputs, examine average performance, or ask whether a model is accurate while leaving the governing architecture unexamined. A polished dossier can conceal unreviewable assumptions. A multidimensional dashboard can hide a rights violation inside an average. A framework can describe monitoring without creating a duty to stop. A sponsor can use an assurance label as a procurement or litigation shield. A shared spreadsheet can display apparently independent ratings without proving that they were deposited before reviewers discussed the case. These are not merely documentation defects; they are failures in the conditions under which a decision claim may be trusted.

Auditable Flourishing addresses a narrower question than “Which theory of flourishing is true?” It asks whether a candidate framework has earned eligibility for a declared comparison and, if several candidates are eligible, what relations among them are supported without compensating away protected floors or forcing a false total ranking. The object of audit is the framework and its claim architecture, not flourishing itself.

This boundary is essential. AF does not certify that a population is flourishing. It does not grant legal conformity, moral adequacy, community legitimacy, procurement approval, safety clearance, or deployment authority. It does not turn plural values into a universal scalar. It creates an inspectable record of what a candidate claims, how it protects non-compensatory conditions, what evidence supports its comparative profile, what remains unknown, and what would require refusal, correction, or retirement.

## 2. Design principles and contribution

The method integrates several established ideas without claiming novelty for them in isolation: rights floors and due process; capability and wellbeing approaches; sustainability and planetary boundaries; program and policy evaluation; educational and psychological measurement; assurance and conformity assessment; algorithmic audit; partial-order and outranking methods; institutional diversity; contestability; and corrigibility. The contribution lies in the object and sequence of evaluation.

First, AF evaluates decision frameworks that seek recurring authority over high-impact flourishing claims. Second, it applies a constitutional-class, non-compensatory eligibility screen before comparative evaluation. Third, it accepts functionally equivalent evidence across computational, legal, institutional, community, and deliberative architectures rather than requiring one favored artifact form. Fourth, it compares eligible candidates through an ordinal partial order without summing category labels. Fifth, it preserves four distinct non-ranking outcomes: scoped dominance, invariant incomparability, non-decisiveness, and refusal. Sixth, it binds every result to version, scope, evidence cut-off, scenario set, host, uncertainty, appeal, expiry, and authorization boundaries.

The protocol is guided by five implementation principles.

1. **Non-compensation.** A material protected-floor or constitutional defect cannot be canceled by aggregate strength elsewhere.
2. **Functional neutrality.** Criteria specify protective functions, not a mandatory software or document form.
3. **Uncertainty preservation.** Indeterminate, not reviewed, non-decisive, and refusal are legitimate outputs rather than errors to be hidden.
4. **Authority separation.** Evaluation, legal conformity, moral adequacy, community legitimacy, and deployment authorization are independent findings issued by different competent processes.
5. **Recursive falsifiability.** AF must expose conditions under which its criteria, anchors, burden, architecture neutrality, reliability, or decision value should be revised, restricted, paused, or retired.

These principles do not establish public legitimacy by authorial assertion. C0–C5 constitute a proposed constitutional class for framework assurance. Until a plural and independent process reaches an applicable constitutional-legitimacy level, positive findings remain protocol-relative.

### 3. Object, scope, and record boundary

A candidate is a versioned decision architecture intended for repeated use. It must identify inputs, transformations, constraints, outputs, governance, refusal behavior, update pathways, and audit artifacts. A list of aspirations is not a framework. A candidate may contain protected implementation details, but the normative and decision core must remain sufficiently inspectable for independent reconstruction, challenge, replay, or rejection.

Every AF administration binds at least the following:

- candidate identity, version, and cryptographic or documentary identifier;
- declared decision scope, population, jurisdiction, and authority claim;
- evidence cut-off and evidence maturity;
- applicable AF profile;
- frozen scenario-set identity, version, and content hash;
- reviewer identities or blinded identifiers, qualifications, conflicts, and assistance disclosures;
- immutable pre-discussion rating deposits;
- affected-party standing, access requirements, and evidence-governance constraints;
- criterion findings, root causes, dissent, adjudication, appeal, and expiry;
- separate legal, moral, community, and deployment findings where supplied by competent authorities.

A result is not portable beyond those bindings. A Stage A finding for one version, scope, or evidence cut-off does not attach to a later release or a wider claim. A Stage B relation established under one scenario set does not automatically generalize to another domain. A copied manuscript or isolated worksheet does not inherit package-level checksum, rendering, or verification attestations.

AF separates four conceptual layers.

| Layer            | Function                                                                                                   | Non-substitution rule                                               |
|------------------|------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------|
| <b>Floors</b>    | Protect dignity, rights, basic needs, due process, non-discrimination, and applicable ecological integrity | Aggregate benefits cannot compensate for a material breach          |
| <b>Fields</b>    | Represent multidimensional improvement above floors                                                        | Averages cannot hide subgroup harm or erase dimensions              |
| <b>Horizons</b>  | Address long-term continuity, tail risk, resilience, and ecological limits                                 | Short-term gains cannot justify ruin-level or irreversible exposure |
| <b>Integrity</b> | Make claims inspectable, challengeable, measurable, updateable, and capable of refusal                     | Documentation without causal governance is insufficient             |

This architecture avoids two opposite errors. A purely aggregative framework may conceal sacrifice inside a high total. A purely veto-based framework may refuse every difficult comparison without showing how admissible alternatives differ. Stage A protects the constitutional boundary; Stage B supports disciplined comparison above that boundary.

## 4. Stage A: non-compensatory eligibility

### 4.1 Six criteria

Stage A asks whether a candidate is eligible for a stated comparison. The six criteria are functional and apply across architecture families.

| Criterion                                | Required function                                                                                                                                 | Illustrative evidence                                                                                                     |
|------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| <b>C0 Inspectability</b>                 | Expose claims, concepts, inputs, transformations, assumptions, scope, authority, and uncertainty sufficiently for independent challenge           | public core; model or rule documentation; authority map; provenance; dependency records; decision procedures              |
| <b>C1 Rights</b>                         | Identify protected floors, affected parties, distribution, non-derogable or non-compensable conditions, redress, and escalation                   | rights register; equality analysis; standing; remedy; legal or community safeguards; emergency limits                     |
| <b>C2 Auditability</b>                   | Preserve a versioned and replayable claim-evidence-decision record                                                                                | evidence graph; logs; scenario versions; configuration; decision notes; dissent; hashes; reproducibility artifacts        |
| <b>C3 Measurement integrity</b>          | Connect indicators and categories to constructs, uncertainty, subgroups, horizons, and decisions without illicit arithmetic                       | construct maps; measurement cards; validity evidence; subgroup analysis; missingness; calibration; anti-Goodhart controls |
| <b>C4 Corrigibility</b>                  | Detect failure and support correction, deprecation, appeal, and retirement                                                                        | monitoring triggers; appeal reversals; sunset clauses; change control; public changelog; reconvening rules                |
| <b>C5 Refusal and execution boundary</b> | Distinguish recommendation from authority and refuse, defer, narrow, or escalate under insufficient evidence, competence, capacity, or permission | abstention rules; jurisdiction boundary; referral; safer alternative; authorization record; shutdown or pause conditions  |

### 4.2 Criterion states

Each criterion receives exactly one state:

- `adequate_for_scope`;
- `adequate_control_dependent`;
- `indeterminate`;
- `material_defect_patchable`;
- `material_defect_unpatchable_for_stated_claim`;
- `not_reviewed`.

The materiality test is causal. A gap is material when a credible correction could change eligibility, protected-floor routing, Stage B applicability or relation, refusal, authority, public claims, or the burden and access feasibility of the review. Cosmetic incompleteness and reviewer preference do not become material defects by assertion.

A control-dependent positive finding must identify the control, owner, verification record, monitoring trigger, expiry, and requalification consequence. An indeterminate finding records what evidence is missing, why it matters, who could obtain it, and whether the uncertainty is safely repairable. Not reviewed is not evidence of adequacy.

### 4.3 Root causes and non-duplication

A single causal defect can create consequences across several criteria. AF assigns one root-cause identifier and one primary criterion owner. Other criteria may cross-reference the consequence, but the same defect cannot become several independent penalties. This preserves visibility without inflating severity through repeated counting.

### 4.4 Overall routing

Stage A is conjunctive and non-compensatory. Its routing can be represented as follows:

1. Invalid or incomplete administration, inadequate authority, infeasible burden/access, unmanaged conflict, or unrecognized controlled values produce `protocol_refusal_to_evaluate`.
2. Any established unpatchable material defect produces `stage_a_non_admissible_unpatchable_for_stated_claim`.
3. Otherwise, any established patchable material defect produces `stage_a_non_admissible_patchable`.
4. Otherwise, an unresolved protected floor, insufficient evidence, unresolved competence, or any indeterminate/not-reviewed criterion produces `indeterminate_insufficient_evidence`.
5. Otherwise, any criterion that is adequate only because named controls are essential produces `stage_a_admissible_scoped_control_dependent`.
6. Only six adequate\_for\_scope findings produce `stage_a_admissible_scoped`.

This precedence preserves established material defects even when another criterion remains uncertain; the criterion vector still publishes every indeterminate or not-reviewed state.

A positive Stage A label is not a free-standing word. It must include the canonical machine token and the exact candidate, version, scope, evidence cut-off, AF version, independent host, conditions, expiry or requalification date, and appeal route. Generic words such as “approved,” “certified,” “safe,” “ethical,” “legitimate,” or “compliant” are prohibited.

### 4.5 Proportional profiles without downward waiver

AF-Lite, AF-Standard, and AF-High-Consequence vary administrative intensity, not constitutional semantics. Routing depends on stakes, rights impact, reversibility, population and ecological scale, novelty, contestation, and authority. A partially completed route remains `profile_undetermined_incomplete`; it cannot default to AF-Lite. High-consequence cases cannot be routed downward because an institution lacks resources. Capacity constraints require support, delay, narrower claims, indeterminacy, or refusal.

## 5. Stage B: ordinal partial-order comparison

### 5.1 Preconditions

Stage B is available only when candidates are Stage-A admissible for a shared scope and evidence cut-off and when the administration can establish:

- a common set of applicable dimensions;
- no unresolved protected-floor defect;
- a frozen, versioned, matched scenario suite;
- independent and qualified pre-discussion ratings;

- evidence-supported categories and uncertainty;
- documented applicability, dissent, and perturbation alternatives.

A Stage B category is **floor-linked unreachable** when its descriptor necessarily entails a material Stage A defect under the same scope and evidence cut-off. Category 0 on Rights robustness and Refusal quality is therefore not a stable comparative weakness; it routes back to Stage A. Category 1 on either dimension requires a written explanation of why the weakness does not meet the Stage A materiality threshold.

## 5.2 Provisional dimension set

The provisional AF-SB12-v5.6 dimension set contains exactly the twelve ordered dimensions controlled by the v6.2 Core Protocol and Appendix F. Version 6.2 leaves the dimension set unchanged and preserves the separately versioned symmetric relation rule, AF-SB-RELATION-v6.0.

1. Rights robustness
2. Audit completeness
3. Measurement validity
4. Uncertainty discipline
5. Corrigibility
6. Anti-gaming resilience
7. Non-domination
8. Refusal quality
9. Tail-risk adequacy
10. Ecological adequacy
11. Institutional compatibility
12. Reproducibility and public challenge

Each dimension uses ordered categories 0–4, but the numbers are labels for ordered anchors, not cardinal quantities. Evidence maturity, reviewer confidence, applicability, burden, and disagreement are separate metadata fields; they are never substituted for an AF-SB12 dimension or added to an ordinal category. Reviewers may not sum, average, subtract, weight, or convert categories into a scalar “flourishing score.” Indicator-level quantities may be compared only when meaningful units, construct validity, comparable scope, uncertainty, and missingness permit it.

## 5.3 Base relation

For candidates A and B evaluated on the same applicable dimensions, A dominates B when A is no worse on every dimension and strictly better on at least one. B dominates A under the reverse condition. When each is better on at least one dimension, their profiles cross and are incomparable. When all categories are equal, the base relation is a tie.

The base relation is not yet the public outcome. Category anchors may be uncertain, and a qualified reviewer may defend an adjacent category supported by evidence. The protocol must therefore replay the relation under every evidence-supported substitution.

## 5.4 Symmetric perturbation and relation set

Let  $R^*$  be the set containing the base relation and every relation obtained from a valid defended substitution. The public outcome is:

- dominance with `dominates_a` when  $R^* = \{\text{dominates\_a}\}$ ;

- dominance with `dominates_b` when  $R^* = \{\text{dominates\_b}\}$ ;
- incomparable when  $R^* = \{\text{incomparable}\}$ ;
- `non_decisive` when  $R^*$  contains more than one relation, contains a tie, or remains unresolved under valid but potentially repairable conditions;
- `refusal_to_compare` when prerequisites, common scope, evidence integrity, floor status, scenario governance, competence, or authority are invalid.

This rule is symmetric. A fragile dominance relation cannot survive merely because the perturbation test was ignored. Conversely, a stable crossing cannot be erased by peripheral disagreement. If candidates remain crossing under all defended alternatives, the correct result is incomparability, not non-decisiveness.

## 5.5 Worked relation examples

**Invariant dominance.** Suppose candidates are tied on eleven dimensions. On Measurement validity,  $A = 4$  and  $B = 2$ . A qualified reviewer defends  $B = 3$ . The base relation is  $A$  dominates  $B$ ; after substitution,  $A$  still dominates  $B$ .  $R^* = \{\text{dominates\_a}\}$ ; dominance is warranted for the declared scope.

**Invariant crossing.** Suppose  $A$  is stronger on Rights robustness and  $B$  is stronger on Refusal quality. A disputed change to Implementation feasibility does not remove either strict difference. Every relation remains crossing.  $R^* = \{\text{incomparable}\}$ ; the stable information is that neither candidate is globally no worse.

**Non-decisive relation.** Suppose  $A$  is strictly better only because Measurement validity is 4 rather than 3. A qualified alternative places  $A$  at 3, creating a tie.  $R^* = \{\text{dominates\_a}, \text{tie}\}$ ; a dominance claim would overstate the record. The correct result is non-decisive.

These examples show logical behavior, not empirical yield. The frequency of each outcome depends on candidate selection, anchor use, category dependence, reviewer disagreement, scenario difficulty, and evidence quality.

## 6. Administration and independence

### 6.1 Host and sponsor separation

Independence is a mechanism, not a label. Every administration publishes a dossier identifying the sponsor, host, data controller, publication controller, reviewer-selection process, conflict adjudicator, financial and intellectual interests, affected-party representation, host interest in the outcome, host precommitment, and AI-assistance policy.

A claim of independence is invalid when a sponsor or candidate controls data, reviewer selection, adjudication, or publication; when dissent or null and adverse publication rights are absent; when the host has a predetermined outcome; or when shared model-generated ratings are counted as independent human judgments. Funding can be managed through declared controls, but outcome-contingent payment or publication veto is incompatible with an independent challenge.

### 6.2 Immutable pre-discussion ratings

A workbook cannot prove that a rating was made before discussion or remained unedited. AF therefore treats the workbook as an aggregation instrument only. Each primary rating must be deposited into an independently controlled system that produces:

- a deposit identifier;
- reviewer, case, decision layer, and target identifiers;
- timestamp before panel discussion;
- content hash;

- signature, append-only audit event, or equivalent proof;
- repository or system identity;
- supersession and withdrawal lineage;
- AI-assistance disclosure.

Retrieval, translation, or formatting assistance may be allowed under a published policy. A model-generated primary rating or shared-model harmonization creates correlated errors and is not treated as an independent human rating for reliability estimation.

### 6.3 Reviewer qualification and halo control

Reviewers complete at least five practice cases, criterion-specific competence checks, Stage B training, independence and conflict training, evidence-equity and accessibility training, and AI-boundary training. Qualification has scope and expiry. Material anchor or protocol changes trigger requalification.

Criterion order is randomized where feasible, and reviewers record a global-impression probe before or separately from detailed ratings. Reliability is reported by criterion, Stage A final outcome, Stage B dimension, and Stage B final relation. These layers cannot be blended into one coefficient. Gwet's AC1 or AC2, weighted agreement, Krippendorff's alpha, or related estimators may be appropriate depending on scale and design, but the coefficient does not substitute for disagreement analysis, confidence intervals, prevalence, subgroup effects, or consequences.

### 6.4 Affected-party standing and evidence equity

Affected people and communities may submit evidence, challenge scope and material facts, request accessible reasons, receive interpretation or disability support, preserve dissent, appeal findings, and seek correction or withdrawal of improperly obtained evidence. Non-documentary, oral, customary, observed-practice, translated, and community-governed evidence may be functionally adequate when provenance, consent, purpose limitation, custody, interpretation, challengeability, compensation, withdrawal conditions, uncertainty, and data sovereignty are recorded.

AF must not confuse documentation polish with safeguard quality. Assisted dossier pathways can help low-resource candidates make existing safeguards inspectable, but assistance cannot invent a safeguard or change an adverse finding for cosmetic reasons.

### 6.5 Burden and access as gates

Feasibility is conjunctive. A review is not feasible unless reviewer hours, adjudication hours, calendar time, qualified reviewer count, and required access support are all sufficient. Favorable per-case averages cannot create a false green result when calendar capacity, reviewer availability, interpretation, or accessibility fails.

Burden is reported separately for completed feasible administrations and infeasible, stopped, or aborted attempts. The record includes preparation time, review time, adjudication, translation, access support, direct cost, uncompensated labor, delay, withdrawal, abandonment, and subjective burden by role. A protocol that detects defects but imposes disproportionate or unequal burden may require redesign or retirement.

## 7. Authority and anti-shield boundary

AF uses an authority-specific verdict vector with five independent components:

1. AF eligibility;
2. legal conformity;
3. moral adequacy;
4. community legitimacy;

## 5. deployment authorization.

The AF host may issue only the first. A competent jurisdiction-specific authority supplies legal findings; legitimate moral or community processes supply their findings; and an accountable operating authority supplies deployment authorization. One component does not entail another.

An AF finding cannot operate as a certification, due-diligence safe harbor, liability defense, warranty, procurement approval, legal opinion, or conclusive proof of reasonable care. Procurement documents, board papers, regulatory submissions, or litigation materials that mention AF must reproduce the exact bounded status, token, scope, evidence cut-off, host, conditions, dissent, expiry, known limitations, appeal, and verdict vector. AF does not replace domain-specific law, professional duty, human-rights due diligence, impact assessment, remedy, or continuing oversight.

Legal and rights-sensitive administration must also recognize jurisdictional differences. International and regional derogation regimes, margin-of-appreciation doctrine, non-derogable protections, non-refoulement, minimum core obligations, and business human-rights responsibilities require competent legal interpretation. AF can route and preserve those issues; it cannot decide them as law.

## 8. Scenario governance

Stage B and external replication require more than a scenario form. A scenario set is a governed research object with an identifier, semantic version, content hash, frozen timestamp, domain, intended function, difficulty status, architecture-equivalence criteria, translations and accessibility materials, leakage rules, change control, and retirement conditions.

Architecture equivalence does not require identical artifacts. It requires comparable decision stakes, evidence maturity and missingness, time budget, remedy access, and protected-floor exposure. Scenario revisions create a new version and hash. A scenario is retired when it leaks, ceases to discriminate, embeds unjustified architecture assumptions, creates disproportionate harm, or no longer represents the intended construct.

Two hosts cannot claim replication if they silently use different scenario content or difficulty. Every comparison and rating deposit therefore binds the scenario-set version and hash.

## 9. Phase 0 and external validation

A full external programme may require multiple hosts, diverse candidate families, affected-party participation, architecture-fairness analysis, multiple comparator arms, ablation, burden accounting, and downstream outcome tracking. Attempting every component at once risks making the protocol unexecutable. Version 6.2 therefore defines a minimum viable Phase 0.

Phase 0 uses:

- one independent host;
- twenty frozen synthetic cases;
- four to six trained reviewers;
- immutable pre-discussion deposits;
- randomized criterion order and global-impression probes;
- criterion- and layer-specific reliability reporting;
- AF v6.2 and an audit-theatre decoy, with a structured checklist where feasible;
- exploratory analysis only;
- no live affected-party participation;
- no architecture-fairness, field-validity, or deployment claim.

The primary questions are whether the administration can be executed as specified, whether blank and partial records fail closed, whether reviewers can use the anchors, whether ratings are deposited independently, whether the protocol distinguishes causal safeguards from polished documentation, and what burden and adverse events arise. Twenty cases and four to six reviewers are not treated as a confirmatory sample. They provide prevalence, variance, disagreement, missingness, and burden estimates for simulation-based design of a later study.

A Phase 0 budget separates reviewer compensation, host administration, access and translation, repository and data costs, and contingency. Progression requires data-integrity and independence compliance, administrable burden, informative criterion-level estimates, audit-theatre discrimination, and unconditional publication of null and adverse findings. Failure can justify revision, restriction, pause, or retirement rather than automatic expansion.

For later validation, strong comparators include a simpler structured checklist, domain standard of care, a minimal-documentation control, and an audit-theatre decoy. Fractional or prioritized ablation is preferable to an exhaustive one-component-at-a-time design when the full factorial programme would be infeasible. Outcomes include false clearances, false rejections, uncertainty reduction, decision changes, delay, access failures, appeal reversals, remedy quality, affected-party value, architecture-linked differential error, and downstream effects.

A Registered Report is a strategically suitable publication form because Stage 1 peer review can evaluate the question, design, estimands, and analysis before results are known. Preparing a Stage 1 protocol does not imply journal acceptance. The release therefore includes a submission-ready outline and deposit metadata but makes no acceptance or DOI claim.

## 10. Structural yield and falsification

Under independent uniform five-category dimensions, the probability that one random profile dominates another across  $n$  dimensions is:

$$P(\text{dominance}) = 2[(3/5)^n - (1/5)^n].$$

At  $n = 12$ , the result is approximately 0.4354%. If all twelve dimensions are restricted to categories {2,3}, the corresponding independent probability is approximately 6.2864%. If five dimensions are binary-restricted and seven retain five categories, the result is approximately 1.3285%. These are structural sensitivities, not forecasts. Stage A selection may restrict the category range, but real outcome frequencies also depend on dependence, scenario selection, reviewer disagreement, applicability, and perturbation structure.

The protocol therefore preregisters competing hypotheses rather than asserting that dominance, incomparability, or non-decisiveness must prevail. High non-decisiveness may honestly represent underdetermination, or it may indicate anchors too coarse for the selected population. Persistent empty comparison sets, low decision value, architecture-linked exclusion, unreliable anchors, inability to detect audit theatre, or burden that exceeds benefit are falsification or revision signals.

AF should be revised, restricted, paused, or retired when independent evidence shows, among other conditions:

- C0–C5 systematically favor an affiliated or document-rich architecture without functional justification;
- an architecturally distinct adequate framework is wrongly excluded, or an affiliated entrant systematically benefits from native artifact resemblance;
- criterion or relation reliability is too low for the intended inference;
- ratings are driven by halo effects or shared AI outputs;
- burden, access failure, or retaliation is disproportionate;
- positive findings are repeatedly misused as legal or procurement shields;
- Stage B produces relations that do not improve decisions or defect localization;
- a simpler method achieves equal or better protection and decision value at materially lower burden.

## 11. Limitations

The principal limitations are empirical and constitutional rather than hidden. AF has not demonstrated inter-reviewer reliability, construct validity, architecture fairness, affected-party legitimacy, predictive validity, causal decision improvement, downstream benefit, or safe live administration. C0–C5 have a published genealogy, rivals, boundary cases, dissent route, and a graduated legitimacy ladder, but no document can establish universal legitimacy. The current scenario library and reviewer curriculum are specifications awaiting independent use. The provisional Stage B dimensions may prove redundant, incomplete, too coarse, or burdensome. The reference implementation tests internal logic, not institutional behavior. A hash proves artifact identity, not truth.

These limitations determine the correct readiness claim. AF v6.2 is internally specification-complete, synthetic-replay-ready, reviewer-training-ready, and independently governed Phase 0 design-ready. It is not an empirically validated standard or live high-consequence authority.

## 12. Conclusion

Auditable Flourishing proposes a disciplined answer to a recurring governance problem: before a framework can guide high-impact flourishing claims, it should show that its architecture is inspectable, rights-protecting, auditable, measurable, corrigible, and capable of refusal. Those conditions cannot be averaged away. Frameworks that meet them can be compared without forcing plural values into a scalar rank. Frameworks that do not meet them should be patched, narrowed, held indeterminate, or refused.

Version 6.2 moves the programme from a strong internal specification toward an executable independent test. It corrects artifact drift, fail-open workbook defaults, blended reliability, missing scenario ownership, absent immutable deposits, underowned reviewer training, unbounded legal characterization, and an overambitious pilot path. The result is not closure of the empirical question. It is a sharper instrument for allowing evidence, affected parties, independent reviewers, and real institutional constraints to determine whether the method deserves to persist.

## AF v6.2 synchronization note

This artifact implements the v6.2 three-axis record (decision\_status, evidence\_state, review\_state), the non-erasure invariant, perturbation-replay anti-inflation and anti-suppression controls, and the separation of raw agreement, reliability, adjudication, and external validity. These controls clarify the existing protocol and do not confer legal, moral, procurement, or deployment authority.

## References

- American Educational Research Association, American Psychological Association, & National Council on Measurement in Education. (2014). *Standards for educational and psychological testing*. American Educational Research Association.
- Campbell, D. T. (1979). Assessing the impact of planned social change. *Evaluation and Program Planning*, 2(1), 67–90. [https://doi.org/10.1016/0149-7189\(79\)90048-X](https://doi.org/10.1016/0149-7189(79)90048-X)
- Center for Open Science. (n.d.). *Registered Reports*. <https://www.cos.io/initiatives/registered-reports>
- Council of Europe. (1950). *Convention for the Protection of Human Rights and Fundamental Freedoms*, Article 15.
- European Court of Human Rights. (2025). *Guide on Article 15 of the European Convention on Human Rights: Derogation in time of emergency*.
- Gwet, K. L. (2014). *Handbook of inter-rater reliability* (4th ed.). Advanced Analytics.

International Auditing and Assurance Standards Board. (2013). *International framework for assurance engagements*.

International Organization for Standardization. (2020). *ISO/IEC 17000:2020 Conformity assessment: Vocabulary and general principles*.

Joint Committee on Standards for Educational Evaluation. (2011). *The program evaluation standards* (3rd ed.). SAGE.

Krippendorff, K. (2018). *Content analysis: An introduction to its methodology* (4th ed.). SAGE.

Messick, S. (1995). Validity of psychological assessment. *American Psychologist*, 50(9), 741–749.  
<https://doi.org/10.1037/0003-066X.50.9.741>

Organisation for Economic Co-operation and Development. (2019). *Better criteria for better evaluation: Revised evaluation criteria definitions and principles for use*.

Ostrom, E. (1990). *Governing the commons*. Cambridge University Press.

Raji, I. D., Smart, A., White, R. N., et al. (2020). Closing the AI accountability gap. *Proceedings of FAccT 2020*, 33–44. <https://doi.org/10.1145/3351095.3372873>

Roy, B. (1968). Classement et choix en présence de points de vue multiples. *Revue française d'informatique et de recherche opérationnelle*, 2(V1), 57–75.

Sen, A. (2017). Reason and justice: The optimal and the maximal. *Philosophy*, 92(1), 5–19.  
<https://doi.org/10.1017/S0031819116000309>

United Nations General Assembly. (1966). *International Covenant on Civil and Political Rights*.

United Nations Human Rights Office of the High Commissioner. (2011). *Guiding Principles on Business and Human Rights*.