

Auditable Flourishing v6.2 Phase 0 Feasibility and Registered Report Protocol

Status and purpose

Phase 0 is the minimum executable doorway from specification to independent evidence. It is deliberately smaller than the full validation programme. It tests whether reviewers can be trained, records can be deposited immutably, cases can be rated, disagreements can be replayed, audit theatre can be detected, and burden is feasible. It cannot establish live-domain validity, constitutional legitimacy, architecture fairness across societies, affected-party legitimacy, downstream benefit, or deployment authorization.

1. Minimum design

- one host independent of the author and candidate;
- 20 frozen synthetic cases;
- four to six reviewers; reviewer count supports disagreement analysis but does not replace case count for precision;
- current Reviewer Training and Qualification Standard;
- current scenario-set version and holdout controls;
- immutable pre-discussion deposit system;
- randomized or balanced criterion order;
- global-impression probes;
- Stage A C0-C5 ratings and final outcome;
- selected Stage B base and perturbation cases;
- audit-theatre comparator;
- full burden, access, error, and administration log;
- unconditional publication of null, adverse, and protocol-failure results.

No live affected-party evidence is collected. No statement about affected-party legitimacy is permitted.

1A. Materials-readiness gate

Object	Acceptance evidence
Twenty-case library	frozen version, hashes, equivalence/difficulty review, holdout/leakage plan
Reviewer pool	four to six qualified reviewers with practice and conflicts
Deposit system	tested trusted time, hashes, append-only/supersession audit and blinded access
Perturbation register	admission, adjudication, appeal and count workflow tested
Independent host	signed role, interest, custody and unconditional-publication dossier

Object	Acceptance evidence
Analysis package	versioned scripts, clustered intervals, exclusions and null/adverse reporting
Capacity plan	host-specific hours, calendar, costs, accessibility and contingency assumptions

Included records are templates and examples, not a completed calibrated library or qualified pool.

2. Exploratory questions

Phase 0 estimates:

- completion and missingness by criterion and reviewer;
- criterion and final-layer agreement;
- reasoning and root-cause errors;
- rights/refusal safety misses;
- base versus invariant Stage B relations;
- criterion-order and global-impression effects;
- audit-theatre false clearance;
- reviewer and administration time;
- direct and uncompensated cost;
- rating-deposit failures;
- scenario ambiguity and leakage;
- AI-assistance effects where included.

3. No confirmatory overclaim

The design does not claim that 20 cases and four to six reviewers can verify a universal reliability threshold. At an illustrative agreement proportion of 0.80, twenty independent cases yield an approximate 95 percent interval of 0.62 to 0.98, a half-width of about 0.18 before accounting for clustering, prevalence concentration, missingness, or model choice. More reviewers increase the number and structure of pairwise observations; they do not increase the number of independent cases. Cases buy precision, while reviewers buy richer disagreement structure.

Coefficients and intervals are therefore descriptive and exploratory. With six possible final Stage A outcomes and only twenty cases, final-status agreement is reported with exact counts and raw agreement; a chance-adjusted final-status coefficient is not treated as stable unless the preregistered analysis establishes adequate support. If resources permit, an independent host may raise the frozen case set to 30–40 before registration. Results select a larger design through simulation-based precision or decision-error planning. If the estimates are too unstable to plan, the correct outcome is indeterminate or protocol revision.

4. Fractional ablation ladder

Phase 0 does not run every possible one-at-a-time ablation. It prioritizes:

1. removing or weakening the burden gate;
2. removing immutable deposits;

3. blending criterion and final-layer reliability;
4. removing C1 or C5 routing in controlled synthetic cases;
5. removing perturbation replay;
6. comparing full AF with a simpler checklist and audit-theatre decoy.

Later phases use grouped or fractional designs for additional criteria, dimensions, records, and procedures. The selected design and estimand are fixed before results.

5. Capacity and budget model

Feasibility is computed from the maximum of resource bottlenecks, not from a favorable average. The budget contains:

Fixed resources

- host governance and conflict process;
- protocol adaptation and preregistration;
- scenario preparation and independent content review;
- reviewer training and qualification;
- deposit and data infrastructure;
- statistical analysis plan;
- ethics, privacy, security, and legal review as required;
- publication and archiving.

Variable resources

- reviewer case hours;
- adjudication and appeal hours;
- candidate or dossier support where applicable;
- translation and accessibility;
- data management and security;
- independent statistician and quality assurance;
- direct payments and uncompensated time.

The formula-based budget template records low, central, and high assumptions. The study cannot be declared feasible unless reviewer hours, adjudication hours, qualified reviewer count, calendar, access support, infrastructure, and unconditional publication control are all adequate.

6. Progression gates

Phase 1 may proceed only if:

- no critical unresolved rights, data, or independence event remains;
- deposit integrity is demonstrated;
- reviewers complete the task with interpretable reasons;
- scenario ambiguity and protocol failures are bounded enough to design the next study;
- burden is feasible with funded support;
- audit-theatre false clearance is not unacceptably high;

- an independent body approves the next protocol and sample-size design.

Possible outcomes are retain, revise, restrict, pause, retire, or indeterminate.

7. Registered Report preparation

The release includes a Stage 1 outline containing:

- research questions and competing hypotheses;
- protocol and version freeze;
- sampling and exclusions;
- scenario and reviewer qualification;
- immutable deposit and blinding;
- outcomes and coefficients;
- analysis and uncertainty;
- stopping and adverse events;
- deviations and amendments;
- data, code, and publication plan;
- decision rule.

A Registered Report is preferred because review before results and in-principle acceptance can reduce outcome-dependent publication. It is not mandatory, and no acceptance is claimed. Where no Registered Report is used, equivalent preregistration and publication firewalls remain binding.

8. Phase 0 public report

The report publishes:

- exact files, versions, hashes, and deviations;
- host, sponsor, conflicts, and interests;
- reviewer qualification and AI-assistance classes;
- deposit-system integrity;
- case and criterion completion;
- criterion and final-layer results separately;
- base and perturbation relation sets;
- audit-theatre outcomes;
- burden and cost cohorts;
- adverse and protocol-failure events;
- limitations and next-step decision.

Cited works

Center for Open Science. (n.d.). Registered Reports. <https://www.cos.io/initiatives/registered-reports>

Royal Society Open Science. (n.d.). Registered Reports. <https://royalsocietypublishing.org/rsos/registered-reports>