

Auditable Flourishing v6.2 Reviewer Training, Qualification, and Rating-Deposit Standard

Status and role

This standard owns reviewer preparation, qualification, maintenance, immutable pre-discussion deposits, criterion-order control, global-impression probes, and AI-assistance classification. Completion of training does not make a reviewer competent for every domain, law, community, or profile.

1. Competence model

Reviewer qualification has five components:

1. protocol semantics and non-claims;
2. criterion and dimension interpretation;
3. rights, refusal, authority, and materiality routing;
4. evidence, uncertainty, root cause, and perturbation replay;
5. domain, legal, technical, community, or lived-experience competence needed for the case.

The first four can be trained by the AF programme. The fifth requires independent evidence and may require panel composition rather than one reviewer.

2. Curriculum

The core curriculum covers:

- object of audit and scope binding;
- C0-C5 states and non-compensation;
- material defect versus ordinary weakness versus indeterminacy;
- affected-party standing and evidence equity;
- root-cause non-duplication;
- Stage B anchors and no ordinal arithmetic;
- floor-linked category re-entry;
- symmetric perturbation and relation-set outcomes;
- authority-specific verdict vector;
- burden, access, independence, anti-shield, and host-bias controls;
- scenario binding and leakage;
- AI assistance and protected data;
- public reporting, appeal, expiry, and requalification.

3. Practice and qualification bank

The public practice bank contains at least eight cases covering:

- one clear adequate case;
- one patchable material defect;
- one unpatchable-for-stated-claim defect;

- one indeterminate case;
- one protocol-refusal case;
- one audit-theatre case;
- one invariant dominance perturbation case;
- one invariant crossing and one unstable relation case.

At least two cases contain overlapping symptoms from one root cause. At least two use non-documentary or low-resource evidence. At least one tests host or sponsor pressure.

Qualification is based on safe administration and rationale quality, not forced agreement with an answer key. A reviewer must:

- complete required fields;
- avoid any critical C1/C5 routing error on final qualification cases;
- distinguish candidate defect from protocol or evidence failure;
- provide evidence-linked and falsifiable reasons;
- apply root-cause and perturbation rules correctly;
- disclose uncertainty, conflicts, and AI assistance;
- meet any domain-specific competence requirement.

Failure leads to retraining, narrower role, supervised practice, or non-qualification. A disagreement supported by evidence is not automatically failure.

4. Maintenance and requalification

Qualification expires after a declared interval, material protocol or anchor change, long inactivity, repeated administration error, undisclosed conflict or AI use, or significant domain change. Requalification uses a holdout case and targeted modules. Training and qualification records are versioned and public or securely auditable.

5. Immutable rating-deposit standard

A deposit record contains:

- deposit, study, case, reviewer, layer, criterion or dimension, and candidate identifiers;
- protocol, anchor, scenario, and codebook versions;
- rating, applicability, confidence, evidence pointers, rationale, and uncertainty;
- submission time from the independent system;
- content hash;
- signature or audit proof;
- repository or data-capture system;
- whether the record existed before discussion;
- AI-assistance disclosure;
- withdrawal, correction, or supersession linkage.

The host closes deposits before revealing ratings. Corrections do not overwrite the original; they create a superseding record with reason and time.

6. Criterion order and global impression

Each case uses a preregistered balanced or randomized criterion order. The reviewer records a non-dispositive global-impression probe before criterion work and after completion. The host analyzes order and halo effects separately from criterion reliability.

Reviewers must not discuss candidate quality, share ratings, or use a common model to generate ratings before deposit closure.

7. AI-assistance classes

Class	Examples	Reliability treatment
A0 None	no AI tool used	eligible if other conditions hold
A1 Clerical	spelling, formatting, file navigation	eligible with disclosure
A2 Retrieval/translation	locating text, translation, summarization without rating recommendation	eligible only with verification, data authority, and disclosure
A3 Interpretive	analysis suggestions, anchor interpretation, evidence synthesis	separate sensitivity condition; not presumed independent
A4 Rating-suggestive	proposed category or outcome	excluded from independent human reliability
A5 Rating-generative or harmonizing	model generates or reconciles ratings across reviewers	model-assisted experimental arm only; correlated ratings cannot represent independent humans

A shared model or shared prompt can create common-method correlation even when reviewers edit outputs. Tool, model or version, instructions, evidence sent, output retained, verification, and influence are recorded.

7A. Perturbation-proposal competence

Qualification includes drafting evidence-specific alternatives, identifying relation materiality, distinguishing single from coupled proposals, detecting duplicates and strategic inflation, applying independent admission/adjudication, and preserving supported dissent. Repeated unsupported proposal or suppression requires remediation or requalification.

8. Phase 0 reliability purpose

Phase 0 uses four to six trained reviewers and 20 frozen synthetic cases to estimate prevalence, completion time, missingness, order effects, disagreement patterns, and plausible uncertainty. It is exploratory. Point thresholds and lower confidence bounds are not confirmatory at this scale.

A confirmatory study uses Phase 0 estimates to simulate case and reviewer counts for a prespecified interval-width, classification-error, or decision-error objective. The study reports each criterion and decision layer separately and directly estimates the final conjunction.

9. Reviewer welfare and burden

Reviewers can recuse for conflict, insufficient competence, distress, protected-knowledge restrictions, or unreasonable burden. Recusal is not a candidate defect. Work involving traumatic evidence includes warning, support, exposure limits, and compensation. Uncompensated expert or lived-experience labor is reported.

10. Acceptance tests

A human-independent-reliability claim fails when:

- deposits are absent, mutable, late, or visible before closure;
- reviewers discuss before deposit;
- criterion order or global-impression handling is undisclosed;
- AI-generated or harmonized ratings are counted as independent humans;
- reviewer competence is unbound to the case;
- final outcome agreement is used to conceal weak criterion agreement;
- unlike ordinal and nominal layers are blended;
- the sample is described as confirmatory without a defensible precision design.

Cited works

Gwet, K. L. (2014). *Handbook of inter-rater reliability* (4th ed.). Advanced Analytics.

Harris, P. A., Taylor, R., Thielke, R., Payne, J., Gonzalez, N., & Conde, J. G. (2009). Research electronic data capture (REDCap): A metadata-driven methodology and workflow process for providing translational research informatics support. *Journal of Biomedical Informatics*, 42(2), 377-381. <https://doi.org/10.1016/j.jbi.2008.08.010>

Krippendorff, K. (2018). *Content analysis: An introduction to its methodology* (4th ed.). SAGE.