

Auditable Flourishing v6.2 Scenario Governance and Calibration Standard

Status and role

This standard owns the scenario function. AF v6.2 supplies the design and examples; it does not claim that an independent frozen calibrated twenty-case library exists required by Stage A calibration, Stage B matched comparison, reviewer training, Phase 0, and external replication. A scenario form alone is not a scenario-governance system. Every claimed matched scenario suite must resolve to a frozen version, content hash, calibration record, equivalence rationale, access rule, and retirement history.

1. Scenario objects

The system distinguishes:

- scenario set;
- individual scenario;
- evidence packet;
- hidden answer key or adjudication reference where defensible;
- candidate-specific interface or translation;
- reviewer instructions;
- expected functions and failure modes;
- calibration and difficulty data;
- release and retirement record.

Each object has a stable identifier, semantic version, creation and freeze date, SHA-256, owner, access class, licence or rights note, and supersession link.

2. Required scenario-set record

A set record includes:

- declared domain and decision function;
- intended AF profile and evidence maturity;
- rights, tail-risk, reversibility, and affected-party features;
- architecture-equivalence criteria;
- evidence volume and documentation-intensity targets;
- difficulty and ambiguity targets;
- expected candidate and protocol defects;
- protected or community-governed content rules;
- language and accessibility variants;
- holdout status and leakage controls;
- calibration sample and results;
- change, deprecation, and retirement rules.

3. Frozen matched comparison

Candidate comparisons use the same governing scenario-set version and substantively equivalent evidence conditions. Equivalence does not require identical artifact form. It requires that differences in interface, language, format, or access do not change the decision function, stakes, available evidence, time, remedy, or protected-floor exposure without being declared and analyzed.

Any post-outcome change to scenario content, scoring anchors, evidence packet, or difficulty invalidates the original comparison unless processed as a versioned amendment and full replay.

4. Calibration

Calibration tests:

- whether qualified reviewers recognize the intended function and material defects;
- whether scenarios span clear pass, clear defect, indeterminate, boundary, audit-theatre, perturbation, refusal, and access cases;
- whether difficulty is excessively easy, impossible, or architecture-specific;
- whether evidence supports more than one defensible category;
- whether translations and accessibility variants preserve the governing construct;
- whether hidden answer keys are justified or merely encode author preference.

Scenario calibration is distinct from candidate evaluation. Calibration results cannot be tuned after candidate identities or outcomes are known.

5. Architecture and evidence-form equivalence

The set must include materially different evidence forms where relevant: computational traces, legal reasons, institutional records, oral or observed processes, community-held evidence, deliberative records, and low-resource dossiers. Reviewers test whether the scenario recognizes functionally equivalent safeguards rather than one preferred artifact.

If an affiliated architecture performs better because the scenario mirrors its native concepts or files, the scenario is biased even if the formal rubric is neutral.

6. Holdouts and leakage

Training, qualification, pilot, and confirmatory sets are separated. Reviewers and candidates are informed which materials are public training cases and which are controlled holdouts. Access is logged. Leakage, unauthorized sharing, prior exposure, or model-training contamination is recorded and can require case replacement or analysis restriction.

Synthetic cases used in Phase 0 are published after deposit closure unless security, rights, or study integrity requires delayed release. A confirmatory study must not reuse cases whose answer pattern became part of reviewer qualification without a declared reason and sensitivity analysis.

7. Scenario change control

Changes are classified:

- editorial: no semantic or difficulty effect;
- operational: evidence format, translation, accessibility, or administration change;
- substantive: decision function, stakes, evidence, rights, ambiguity, answer key, or calibration change;
- urgent: security, rights, corruption, or exposure repair.

Operational and substantive changes create a new version and require equivalence or recalibration. Substantive changes invalidate inherited outcome claims unless replayed.

8. Retirement

A scenario is retired when it is leaked, culturally or legally obsolete, no longer construct-valid, unsafe, systematically architecture-biased, unable to discriminate relevant functions, or excessively burdensome. Retirement preserves the record and reason; it does not erase prior results. Results relying materially on a retired scenario are requalified.

9. Acceptance tests

A scenario set cannot support matched comparison or replication when:

- no exact version or hash is available;
- difficulty or evidence conditions differ materially without analysis;
- equivalence criteria are absent;
- the set was changed after outcomes were inspected;
- answer keys encode disputed normative conclusions as fact;
- protected content lacks authority or data-governance controls;
- leakage is material and unaddressed;
- an architecture-family effect is attributable to scenario form rather than candidate protection.